Yusef Syed

Software Engineer: AI Agents, Security & Developer Tools

yusef.syed@mail.utoronto.ca | github.com/YusefSyed | scopehaven.dev

linkedin.com/in/yusef-syed-4b811a289 | Toronto, ON | U.S. & Canadian citizen, no sponsorship needed

Available Summer 2027 (May–Aug); also open to Jan–Apr 2027 | Open to relocating to San Francisco

Education

University of Toronto **Expected May 2030**

B.Sc., Mathematical & Physical Sciences (first year); intended Math and CS double major *Toronto, ON*

Coursework: Linear Algebra, Proofs & Calculus, Statistics; Data Structures (LCCC, dual enrollment)

Experience

Scopehaven | scopehaven.dev **Oct 2026 – Present**

Founder & Engineer; security testing for AI agents *Toronto, ON*

- Shipped **Scopehaven 0.7.0-alpha** in six days: a Node.js CLI and GitHub Action that tests AI agents for access-control failures by checking what they did, not what they reported (3,800+ tests, zero dependencies).
- Ran an internal pilot on an open-source AI workspace (synthetic data): after a user was suspended, an in-flight agent run still wrote a file in **2 of 14** cases; a partial local patch passed 14/14 on the tested path.
- Red-teamed the product's own checks with 57 bypass attempts, then fixed them until false passes fell from **20 to 0**.
- Sole human reviewer for a fleet of parallel AI coding agents: **530+ PRs merged in 26 days**, each gated by tests, CI, code review and secret scanning before auto-deploy.

Independent Security Research | credited as Yusef Syed (Scopehaven) **Oct 2026 – Present**

Security Researcher *Remote*

- Filed **8 private security advisories** (authorization, credential and data-exposure flaws; two rated High) with 5 open-source projects, each reproduced 3/3 with controls first; all in maintainer triage.
- Confirmed **28 bugs**, including cross-user data overwrite and entitlement bypass, in AI-assisted assessments of two self-built apps (local copies, synthetic data), each reproduced 3/3.

Open Source

14 merged PRs: Microsoft, Apache Arrow, W&B, Inspect Scout, Kornia | Full list **Aug 2026 – Present**

- Microsoft PyRIT (AI red-teaming): added package-hallucination attack techniques for Dart, Perl and Raku.
- W&B RAI Toolkit: landed five evaluation fixes: a concurrency fix the maintainer verified at 3.34s to 1.10s, and keeping failed predictions from scoring as passes.
- Microsoft Agent Lightning: landed four reliability fixes, e.g. case-insensitive auth-header replacement.
- Open under review: PyTorch (autograd gradients), JAX, SciPy, Apache Airflow and Sentry.

Projects

Providence | **1st Place, Aramco Americas Challenge**, Hack the North 2026 | Devpost **Sep 2026**

- Team of four at Canada's largest hackathon. Built the Apple Watch client (Swift/SwiftUI) for the team's computer-control AI agent: spoken requests and approvals, with results verified before success is shown.

Agent Eval Mutation Lab | *Python, Inspect, Docker, SQLite* | Results **Aug 2026 – Present**

- Built an Inspect evaluation that audits agent tool effects in isolated Docker/SQLite sandboxes: all **21 decided verdicts** matched an independent SQLite oracle across 65 evidence views (44 abstained).

Aesthetics AI & **Tiraz** | *TypeScript, React Native, PyTorch* | Study **May 2026 – Present**

- Released two iOS apps on the **App Store**: an AI physique and nutrition coach, and an AI outfit-check closet app.

Technical Skills

Languages: Python, TypeScript, JavaScript, SQL, Swift

AI & security: LLM agents, LangGraph, Inspect, Promptfoo, PyRIT, PyTorch; access-control testing (IDOR, RLS), OWASP LLM/Agentic Top 10

Systems & tools: Node.js, Next.js, React Native, SwiftUI, PostgreSQL, Supabase, Docker, GitHub Actions, Git, Claude Code

Awards & Activities

AP Scholar with Distinction | **USA Powerlifting:** #2 in Pennsylvania, #19 in the U.S. (60 kg, 2025) |
Feeding the Valley, Vice President (2025–26): 60+ volunteers, $20,000 raised, 1,700 meals for food banks