

YUSEF SYED

Founder, Scopehaven | Testing whether AI agents stay inside their permissions

yusef.syed@mail.utoronto.ca | scopehaven.dev | github.com/YusefSyed

[linkedin.com/in/yusef-syed-4b811a289](https://www.linkedin.com/in/yusef-syed-4b811a289) | Toronto, ON | U.S. and Canadian citizen

Open to engineering roles at YC companies and high-growth startups | Open to relocation

EDUCATION

University of Toronto

Expected May 2030

B.Sc., Mathematical & Physical Sciences (First Year)

Toronto, ON

Intended double major: Mathematics and Computer Science | Projected GPA: 4.0/4.0 (first-term grades pending)

EXPERIENCE

Scopehaven | scopehaven.dev

Oct 2026 – Present

Founder

Toronto, ON

- Created a local workbench that tests AI agents by checking what each test user can actually reach and, where it can see the backend, what was recorded; it **retests every fix** and never counts missing evidence as a pass.
- Designed a permission suite (tenant isolation, revoked access, approvals, refunds); on a synthetic test app (suite 2.3.0, **44 checks**), a deliberately vulnerable build fails **26**; all 26 fixes verified, 18 allowed-use controls passing.
- In an internal pilot of Scopehaven's approach on a pinned open-source AI workspace (local model, synthetic data), a suspended test user's new requests were refused, but in **2 of 14** cases an agent run started before the suspension still wrote a file; a partial, local repair then passed all 14 checks, verified on the file-write path only.
- Piloted locally on an owned Next.js app (real webhook route, all 8 migrations, synthetic data): **24/24** passed.
- Added a response-level HTTP connector that **proves each test user's login** before and after its checks.
- Built v0.6.0 with a **CLI release gate** that can run in CI, plus JSON, HTML report and SARIF exports.
- Launched scopehaven.dev with an early-access form; defined a fixed-scope offer for teams shipping AI agents.

OPEN SOURCE

14 merged upstream PRs across six projects | [Full list](#)

Aug 2026 – Present

- [Apache Arrow](#): extended four PyArrow slice APIs to accept Arrow integer scalars, with regression tests.
- [Microsoft Agent Lightning](#): landed four reliability fixes, including preserving timeout reasons across report retries.
- [Microsoft PyRIT](#): extended the package-hallucination red-teaming scenario to Dart, Perl and Raku.
- [W&B RAI Toolkit](#): landed five evaluation fixes, including keeping failed predictions from scoring as passes.

PROJECTS

Agent Eval Mutation Lab | *Python, Inspect, Docker, SQLite* | [Results](#)

Aug 2026 – Present

- Built an Inspect evaluation that audits agent tool effects in isolated Docker/SQLite sandboxes; across 65 related evidence views, all **21 decided verdicts** matched an independent SQLite oracle and 44 abstained.
- Reproduced **13 mock-model scenarios** twice with identical reports and database bytes; 73 tests pass in CI.

Aesthetics AI & Tiraz | *TypeScript, React Native, PyTorch* | [Study](#)

May 2026 – Present

- Released two iOS apps on the **App Store**: an AI physique and nutrition coach, and an AI outfit-check closet app.
- Study: 3 seeded garment-completion encoders beat a 52.7% top-1 baseline by **5.8 to 6.3 pts** with full context.

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, SQL, Swift

Frameworks & data: Node.js, Next.js, React Native, SwiftUI, PyTorch, LangChain, PostgreSQL, Supabase, SQLite

Tools & testing: Git, GitHub Actions, Docker, pytest, Inspect, Promptfoo, Ollama, Xcode

AWARDS & LEADERSHIP

Hack the North 2026 | **Winner**, Aramco Americas: Best Beginner Hack | [Devpost](#)

Sep 2026

- Team of four; built Providence's Apple Watch client (Swift/SwiftUI) to direct AI agents by voice, with approvals.

Feeding the Valley | Vice President

2025 – 2026

- Helped coordinate **60+ volunteers**; the initiative raised **\$20,000** and prepared 1,700 meals for local food banks.