Yusef Syed

Software Engineer: AI Agents, Security & Developer Tools

yusef.syed@mail.utoronto.ca | github.com/YusefSyed | scopehaven.dev

linkedin.com/in/yusef-syed-4b811a289 | Toronto, ON | U.S. & Canadian citizen, no sponsorship needed

Available Jan–Apr 2027 (will take the term off) and Summer 2027 (May–Aug) | Will relocate to San Francisco

Education

University of Toronto **Expected May 2030**

B.Sc., Mathematical & Physical Sciences (first year); intended Math and CS double major *Toronto, ON*

Coursework: Linear Algebra, Proofs & Calculus, Statistics; Data Structures (LCCC, dual enrollment)

Experience

Scopehaven | scopehaven.dev **Oct 2026 – Present**

Founder & Engineer; security testing for AI agents *Toronto, ON*

- Built **Scopehaven**, a Node.js CLI and GitHub Action that tests whether an AI agent touched the wrong account, tenant or refund, decided from the app's own records rather than the agent's report (0.7.0-alpha; 56 integration modules, zero dependencies).
- Designed its evidence format (DSSE/in-toto exports, an RFC 9162 hash-chained ledger, a countersignature with an RFC 3161 timestamp) and a stdlib-Python checker that re-derives verdicts from the kept records: **88 of 88** re-derived on the sample pack (synthetic app, no model), 8 of 8 tampered copies rejected.
- Built the database checks that catch a committed write even if it is undone mid-run.
- Built **16 incident-class checks** from a corpus of public agent-failure reports; on our own test lab each fails 3 of 3 on a build with that failure planted and passes on the clean build. Each release carries an OpenSSH-signed manifest, a CycloneDX SBOM and VEX.
- Ran the merge gate for parallel AI coding agents as the only human on the repo: exact-head CI, a recorded review by a second session and secret scanning; 57 review verdicts asked for changes, 0 git reverts, 1 regression restored in 20 h.

Independent Security Research **Oct 2026 – Present**

Security Researcher *Remote*

- Filed **8 private security advisories** (authorization, credential and data-exposure flaws; one rated High in my report) with 5 open-source projects, each reproduced 3/3 with controls first; all in maintainer triage.
- Ran an internal pilot on an open-source AI workspace (synthetic data): after a user was suspended, an in-flight agent run still wrote a file in **2 of 14** cases; a partial local patch passed 14/14 on the tested path.

Open Source

14 merged PRs: Microsoft, Apache Arrow, W&B, Inspect Scout, Kornia | Full list **2026**

- Microsoft PyRIT (AI red-teaming): added package-hallucination attack techniques for Dart, Perl and Raku.
- W&B RAI Toolkit: landed five evaluation fixes: a concurrency fix the maintainer verified at 3.34s to 1.10s, and keeping failed predictions from scoring as passes.
- Microsoft Agent Lightning: landed four reliability fixes, e.g. case-insensitive auth-header replacement.

Projects

Agent Eval Mutation Lab | *Python, Inspect, Docker, SQLite* | Results **2026**

- Built an Inspect evaluation that audits agent tool effects in isolated Docker/SQLite sandboxes: all **21 decided verdicts** matched an independent SQLite oracle across 65 evidence views (44 abstained).

Providence | **1st Place, Aramco Americas Challenge** (Best Beginner Hack), Hack the North 2026 | Devpost **2026**

- Team of four at Canada's largest hackathon. Built the Apple Watch client (Swift/SwiftUI) for the team's computer-control AI agent: spoken requests and approvals, with results verified before success is shown.

Aesthetics AI & **Tiraz** | *TypeScript, React Native* | two iOS apps on the **App Store** | Study **2026**

Technical Skills & Awards

Languages: Python, TypeScript, JavaScript, SQL, Swift

AI & security: LLM agents, LangGraph, Inspect, Promptfoo, PyRIT, PyTorch; access-control testing (IDOR, RLS), OWASP LLM/Agentic Top 10, MITRE ATLAS, CWE

Systems & tools: Node.js, React Native, SwiftUI, PostgreSQL, Supabase, Docker, GitHub Actions, DSSE/in-toto, SBOM/VEX, Git, Claude Code

Awards: AP Scholar with Distinction | USA Powerlifting: #2 in Pennsylvania, #19 in the U.S. (2025) | Feeding the Valley, Vice President: $20,000 raised for food banks